
Light52 -- free, open source MCS51 compatible CPU core

Revision 3 - October 1, 2012

Core Datasheet

© José A. Ruiz 2012

OVERVIEW

Light52 is a free, small open-source CPU core compatible to the Intel MCS51 architecture.

While the core is within the performance envelope of other free MCS51 cores, the implementation trades area for speed.

This core is smaller than most free and commercial MCS51 cores and its speed is comparable to that of a 6-clocker -- see section 8.

The full original MCS51 instruction set is implemented with the possible exception of the BCD opcodes (DA and XCHG) which are optional.

This core is in a very early stage of development; it is not fully tested and not documented at all except for the code comments and this file.

All the information presented in this datasheet should be considered preliminary.

FEATURES

- 100% binary compatible to MCS51 (except possibly for optional BCD instructions).
- Speed comparable to a 6-clocker.
- Configurable through VHDL generics.
- Smaller than most other cores.
- Clock rate not much worse than other commercial cores.
- Includes 16-bit timer, UART and I/O ports.
- Additional peripherals and SFRs can be added easily.
- 256 bytes of IRAM -- fixed size.
- Fully synthesizable, static synchronous design with positive edge clocking and no internal tri-states.

Light52 lacks some features usually present in other MCS51 cores and has some important limitations:

SHORTCOMINGS

- No access to off-chip memory.
- Strictly Harvard: XDATA and XCODE spaces can't be merged into a Von Neumann architecture.
- From 2 to 8 clocks per instruction.
- Far slower than most commercial cores: performance/area ratio is worse even though area is much smaller.
- No On-Chip Debugging capability.

1.- Pinout

Table 1: **Core Signal Pinout**

Signal	Direction	Description
clk	input	Clock, active on rising edge.
reset	input	Active high synchronous reset.
rx_d	input	RxD input for on-board UART.
tx_d	output	TxD output from on-board UART.
external_irq[7..0]	input	High-level-sensitive interrupt inputs
p0_out[7..0]	output	Port P0 8-bit output.
p1_out[7..0]	output	Port P1 8-bit output.
p2_in[7..0]	input	Port P2 8-bit input.
p3_in[7..0]	input	Port P3 8-bit input.

2.- Functional Description

Since the MCS51 architecture is already well documented elsewhere, this datasheet will only deal with those aspects of the core which depart from the original.

In this version of the core, there is no support for shared XCODE/XDATA memory spaces (the core performs simultaneous accesses to XCODE and XDATA and there is no wait state or access arbitration logic yet). The MCU memory model is therefore strictly Harvard.

The peripherals included in the MCU core are generally not compatible to the MCS51 peripherals and are somewhat less flexible -- the core trades programmability in run-time for configurability in synthesis time. See section 7 below for a detailed description of available peripherals.

Existing MCS51 programs will generally NOT work unmodified on this core -- code needs to be ported to the available peripherals and their SFRs like it needs to be in any other MCS51 derivative.

Interrupt operation is identical to the original, except for SFR register IP: interrupt priorities are fixed to their default values and the IP SFR is unimplemented. Interrupts can be tailored to any specific application by customizing the MCU VHDL source.

3.- Special Function Registers

Table 2 lists the SFRs implemented in the current version of the core.

Table 2: **Light52 Special Function Registers**

Symbol	Description	Direct Address	Bit Address and Symbol								Reset Value
ACC	Accumulator	E0H	E7	E6	E5	E4	E3	E2	E1	E0	00H
B	B register	F0H	F7	F6	F5	F4	F3	F2	F1	F0	00H
DPH	DPTR high	83H									00H
DPL	DPTR low	82H									00H
IE	IQR Enable	A8H	AF	AE	AD	AC	AB	AA	A9	A8	00H
			EA	-	-	ES	-	-	ET0	-	
PSW	Program Status Word	D0H	D7	D6	D5	D4	D3	D2	D1	D0	00H
			CY	AC	F0	RS1	RS0	OV	-	P	
SP	Stack Pointer	81H									07H
P0	Port 0 outp.	80H	87	86	85	84	83	82	81	80	00H
P1	Port 1 outp.	90H	97	96	95	94	93	92	91	90	00H
P2	Port 2 inp.	A0H	A7	A6	A5	A4	A3	A2	A1	A0	
P3	Port 3 inp.	B0H	B7	B6	B5	B4	B3	B2	B1	B0	
T0CON	Timer 0 Control	88H	8F	8E	8D	8C	8B	8A	89	88	00H
			-	-	TOCEN	TOARL	-	-	-	TOIRQ	
T0L		8CH									00H
T0H		8DH									00H
T0CL		8EH									FFH
T0CH		8FH									FFH
SCON	UART Control	98H	9F	9E	9D	9C	9B	9A	99	98	00H
			-	-	RxRdy	TxRdy	-	-	RxIrq	TxIrq	
SBUF	Data Buffer	99H									
SBPL	Baud Rate L	9AH									(*1)
SBPH	Baud Rate H	9BH									(*1)
EXTINT	External IRQ Flags	C0H	C7	C6	C5	C4	C3	C2	C1	C0	00H
			EIRQ7	EIRQ6	EIRQ5	EIRQ4	EIRQ3	EIRQ2	EIRQ1	EIRQ0	

Notes

- 1 Only if generic UART_HARDWIRED is false, and then write only.
Registers SBPL and SBPH are initialized as per generics UART_BAUD_RATE and UART_CLOCK_RATE.

4.- Interrupt Vectors

Interrupt management is identical to the original MCS51. The only difference is that the five available interrupt request inputs are connected to different sources:

Table 3: **Interrupt Vectors and Sources**

IRQ	Source	Vector	Priority	8051 equivalent
0	External IRQ	0003h	(highest)	IE0
1	Timer 0	000Bh		TF0
2	Unassigned	0013h		IE1
3	Unassigned	001Bh		TF1
4	UART	0023h	(lowest)	RI+TI

The only other difference is that interrupt priorities are fixed and can't be changed in run time. SFR IP is not implemented. Since there's full access to the VHDL source of the MCU this is not a major limitation.

Register IE works exactly like in the original MCS51, and so does instruction RETI.

When an interrupt is serviced, its priority level is stored internally. Until instruction RETI is executed, no other interrupts of equal or lower priority will be serviced *even if enabled in register IE*.

Instruction RETI clears the current interrupt priority level register. Immediately after executing RETI, and before executing the next instruction of the main program, any pending interrupts of lower priority will be serviced.

5.- Object Code Initialization

The object code for the MCU application is contained within the MCU module. The XCODE ROM is initialized at synthesis time with the contents of generic **OBJ_CODE**, which is expected to be defined in a package named **obj_code_pkg**.

This package must be generated separately for each project and can be considered part of the program application rather than part of the core source.

The light52 project has adopted the convention that the package **obj_code_pkg** must be defined in a vhdl file placed within the MCS51 program directories -- for this purpose, the **obj_code_pkg** package can be considered as just another object code format.

This way, the object code for different projects using this core (or for the different code samples within this project) can be neatly separated from the core sources.

The project includes a Python script (directory **/tools/build_rom**) which can be used to produce a suitable **obj_code_pkg** package file from an Intel-HEX object file. The code samples in directory **/test** contain usage examples for this script (**makefiles** and/or BAT files **build.bat**).

While the method chosen for object code initialization is clean and vendor-independent, it has a major drawback: The object code must be available at synthesis time, and every time the code changes the synthesis has to be re-run. This may be a big handicap in certain applications.

Subsequent versions of the core may provide the option to use memory initialization files so that the XCODE memory can be initialized post-synthesis.

6.- Configuration Generics

Some of the core features can be configured through VHDL generics:

Table 4: **Core Configuration Generics**

Generic	Default	Description
CODE_ROM_SIZE	1024	Size of XCODE ROM in bytes. Can't be zero.
XDATA_RAM_SIZE	512	Size of XDATA RAM in bytes. Can't be zero.
OBJ_CODE	(dummy)	Object code to be placed on ROM. See previous section.
USE_BRAM_FOR_XRAM (*1)	false	Use extra BRAM as XDATA RAM.
IMPLEMENT_BCD_INSTRUCTIONS (*1)	false	True to implement DA and XCHG, false to execute them as NOPs.
SEQUENTIAL_MULTIPLIER (*1)	false	Use sequential multiplier instead of combinational.
UART_HARDWIRED	true	True to hardwire UART baud rate, false to make it configurable at run time.
UART_BAUD_RATE	19200	Default baud rate for UART.
UART_CLOCK_RATE	50MHz	Clock rate assumed by UART initialization constants.
TIMER0_PRESCALER	50000	Value of Timer0 prescaler.

Notes

- 1 Unimplemented, will cause a synthesis assertion failure if given a non-default value.

At this early stage of development some of these generics do not work, and others are not checked against bounds.

TBD: This datasheet should explain each of the configuration generics.

7.- Peripheral Modules

The MCU core includes a number of peripheral modules. These peripherals have been designed hastily in order to provide a working environment for the CPU -- they do not have their own separate test bench, for example.

The current version of the MCU ships with a simple, hardwired UART, a 16-bit timer and four 8-bit input/output ports.

TBD: Explain SFR bus in MCU entity and how to add new peripherals

7.1.- UART

The light52 UART is a simplified version of the original MCS51 serial port.

Some of the operational parameters of the UART are hardwired and non-configurable in the current version, not even at synthesis time:

1. Number of stop bits hardwired to 1.
2. Parity hardwired to NONE.
3. Number of bits per character hardwired to 8.

Besides, the 9-bit mode of the original MCS51, which its applications in inter-MCU communication, is unimplemented yet.

Serial port interrupts work the same as in the original serial port (same vector IRQ4 and same interrupt enable flag IE.ES).

The UART core has some limited capability to recover from errors, described in the VHDL source file **light52_uart** and very similar to that of the original UART:

1. Error conditions such as bad start and stop bits are detected and cause the UART to discard the received byte and wait for the next start bit.
2. Bit sampling mismatches are detected and the bit values are decided by majority.

A follow-up version of this core will include flags for those detected errors, as well as TX and RX overruns.

Since all operational parameters are hardwired except possibly the baud rate, the UART setup is easy: set the baud rate by writing to registers SBPL and SBPH and enable interrupts by setting flag IE.ES -- the UART can be operated in polling mode too if desired.

Register SCON

This register reflects the status of the serial port:

	7	6	5	4	3	2	1	0
SCON	0	0	RxRdy	TxRdy	0	0	Rxlrq	Txlrq
	h	h	r	r	h	h	w1c	w1c

Bits marked 'h' are hardwired and can't be modified.

Bits marked 'r' are read only; they are set and reset by the core.

Bits marked 'w1c' (write 1 to clear) are set by the core when an interrupt has been triggered and must be cleared by the software by writing a 1.

TxRdy	Ready to transmit. High when there's no transmission in progress. It is cleared when data is written to SBUF and is raised at the same time a TX interrupt is triggered.
RxRdy	Received data ready. High when there's data in the RX buffer. Raised at the same time the RX interrupt is triggered, cleared when SBUF is read.
Rxlrq	RX interrupt pending service. Raised when the RX interrupt is triggered, cleared when 1 is written to it.
Txlrq	TX interrupt pending service. Raised when the TX interrupt is triggered, cleared when 1 is written to it.

When writing to the status/control registers, only flags **Txlrq** and **Rxlrq** are affected, and only when writing a '1' as explained above. All other flags are read-only.

Interrupt flags are triggered at the following moments:

Txlrq	Last clock cycle of the TX stop bit.
Rxlrq	At 11/16ths of the RX stop bit, only if the stop bit is valid.

Register SBUF

This is the read/write buffer of the serial port. It gives the software access to the 1-byte-deep receive and transmit buffers. These buffers work like the original MCS51 serial port.

	7	6	5	4	3	2	1	0
SBUF	UART Tx/Rx buffer register							
reset	x	x	x	x	x	x	x	x
	r/w	r/w	r/w	r/w	r/w	r/w	r/w	r/w

Bits marked 'r/w' can be read and written to by the CPU and can be updated by the core.

SBUF Writing to this register will trigger a serial port transmission unless SCON.TxRdy=0.
Reading this register will give the last byte received by the UART, if any.

Writing to this register will trigger a transmission unless there is already a transmission going on (flag SCON.TxRdy=0). In which case the last write access will be ignored. there is no overrun flag to signal this event; the user must prevent it from happening.

When a byte is received, the core raises flag SCON.RxRdy=1. If a new byte is received before the last one has been read (i.e. with flag SCON.RxRdy=1), the receive buffer register will be overwritten with the new data. Again, there is no indication that this has happened; the user must make sure to prevent these overruns.

Reading from this register when flag SCON.RxRdy=1 will clear the flag and return the last received byte.

Reading from this register when flag SCON.RxRdy=0 will return undefined data (usually the last received byte but this may change in later versions).

Note that reading SBUF does NOT clear the RxIrq flag. The flag must be cleared explicitly.

Registers *SBPH* and *SBPL*

If generic `UART_HARDWIRED` is set to false, then the UART implements these two write-only registers.

These registers should be loaded with the baud period measured in clock cycles -- no prescaling involved:

$$\text{BIT_PERIOD} = \text{UART_CLOCK_RATE} / \text{UART_BAUD_RATE}$$

The bit period register is the 13-bit wide combination of `SBPH` and `SBPL`, with the 3 higher bits of `SBPH` being ignored.

Note that these registers are write only: reading from their addresses will return an indeterminate value (actually, the value of the `SCON` register). This saves logic and is hardly an inconvenience for the programmer, which will seldom have to read these registers.

These registers are loaded at reset with their default value, defined by generics `UART_BAUD_RATE` and `UART_CLOCK_RATE`, according to the same formula above.

When the generic `UART_HARDWIRED` is set to true, these registers are hardwired to their default value and writing to them has no effect.

Note that the UART is totally independent of the timer and indeed of any other module, unlike the original MCS51.

7.2.- Timer 0

Basic timer, not directly compatible to any of the original MCS51 timers. This timer is totally independent of the UART.

This is essentially a reloadable 16-bit up-counter that optionally triggers an interrupt every time the count reaches a certain value.

Timer Registers

The timer includes 3 registers:

1. A configurable prescaler register of up to 31 bits.
2. A 16-bit compare register accessible through T0CL and T0CH.
3. A 16-bit counter register accessible through T0L and T0H.

Reading T0L or T0H will give the value of the timer register. If the registers are read while the count is enabled, the software has to deal with a possibly inconsistent (T0L,T0H) pair and should apply the usual tricks -- majority vote, etc.

The prescaler is reset to 0 when T0CEN=0.

When T0CEN=1 it counts up to (TIMER0_PRESCALER - 1), then rolls over to 0 and the timer register is incremented.

TIMER0_PRESCALER is a VHDL generic configurable at synthesis time.

The compare register is write-only, in order to save logic. Reading T0CH or T0CL will give the value of T0CON instead.

Timer Operation

The counter register is reset to 0 when T0CEN=0. When flag T0CEN is set to 1, the counter starts counting up at a rate of one count every TIMER0_PRESCALER clock cycles.

When counter register = reload register, the following will happen:

- If flag T0ARL is 0 the core will clear flag T0CEN and and raise flag Irq, triggering an interrupt. The counter will overflow to 0000h and stop.
- If flag T0ARL is 1 then flag T0CEN will remain high and flag Irq will be raised, triggering an interrupt. The counter will overflow to 0000h and continue counting.

Register TSTAT

This register reflects the status of the timer:

	7	6	5	4	3	2	1	0
TSTAT	0	0	TOCEN	TOARL	0	0	0	TOIRQ
reset	0	0	0	0	0	0	0	0
	h	h	r/w	r/w	h	h	h	w1c

Bits marked 'h' are hardwired and can't be modified.

Bits marked 'r' are read only; they are set and reset by the core.

Bits marked 'r/w' can be read and written to by the CPU and can be reset by the core.

Bits marked 'w1c' (write 1 to clear) are set by the core when an interrupt has been triggered and must be cleared by the software by writing a 1.

TOCEN

Count ENable.

Must be set to 1 by the CPU to start the counter.

When TOCEN=0 the prescaler and the counter register are reset to 0.

Writing a 1 to TOCEN will start the count up. The counter will increment until it matches the compare register value (if TOARL=1) or until it overflows (if TOARL=0), at which moment the counter register will roll back to zero.

TOARL

Auto ReLoad. Set to 1 to enable compare/autoreload mode.

TOIRQ

Timer interrupt pending service.

Raised when the timer interrupt is triggered, cleared by writing 1 to it.

Registers TOL,T0H

	7	6	5	4	3	2	1	0
TOL	Counter register value, bits 7..0							
reset	0	0	0	0	0	0	0	0
	r/w	r/w	r/w	r/w	r/w	r/w	r/w	r/w

	7	6	5	4	3	2	1	0
T0H	Counter register value, bits 15..8							
reset	0	0	0	0	0	0	0	0
	r/w	r/w	r/w	r/w	r/w	r/w	r/w	r/w

Bits marked 'r/w' can be read and written to by the CPU and can be reset by the core.

T0H:TOL This is the current value of the counter register. Will be reset to zero when TOCEN=0.

Registers TCL,TCH

	7	6	5	4	3	2	1	0
TOCL	Compare register value, bits 7..0							
reset	1	1	1	1	1	1	1	1
	w	w	w	w	w	w	w	w

	7	6	5	4	3	2	1	0
TOCH	Compare register value, bits 15..8							
reset	1	1	1	1	1	1	1	1
	w	w	w	w	w	w	w	w

Bits marked 'w' can be written to by the CPU but reading them will yield an undefined value.

TOCH:TOCL This is the current value of the reload register.

7.3.- Input/Output Ports

The MCU includes 4 8-bit I/O ports. In order to save logic, the ports are hardwired to be either input or output, and are not configurable even at synthesis time -- it is simpler and cheaper to just add or modify whatever port setup is needed in each particular application than trying to provide for all possibilities in advance.

The port SFR addresses are the same as the original P0..P3 port addresses. However, since the ports are strictly input or strictly output, the behavior of the ports is different in a very important way:

The '**Read-Modify-Write**' behavior of the MCS51 is **not implemented**:

- All instructions reading an input port read the pin regardless of addressing mode.
- All instructions reading an output port read the register regardless of addressing mode.

In short, writing to an input port will not have any effect. Reading an output port will access the port output registers and NOT the pins, as stated above.

The input ports are NOT registered or otherwise proof against metastability. A read instruction will read the instantaneous value of the input pin; then the value *will* be registered to internal register T before reaching its destination within the CPU, but there is some amount of logic between the port input and the T register, which means the port inputs have some non-negligible setup time, and not necessarily the same for all lines of the same port.

Subsequent versions of the core will register the input ports to minimize and equalize setup times.

7.4.- External Interrupt Inputs

The MCU has 8 external interrupt inputs which, in the current version of the core, are meant mostly for debugging.

The inputs are registered and level sensitive. As long as input `external_irq[i]` is high, flag `EXTINT[i]` will be high and the interrupt request line `IRQ0` of the CPU will be asserted.

Subsequent versions of the core will use edge triggering and add a mask register.

Register `EXTINT`

This register contains the external interrupt pending flags:

	7	6	5	4	3	2	1	0
EXTINT	EIRQ7	EIRQ6	EIRQ5	EIRQ4	EIRQ3	EIRQ2	EIRQ1	EIRQ0
reset	0	0	0	0	0	0	0	0
	w1c	w1c	w1c	w1c	w1c	w1c	w1c	w1c

Bits marked 'w1c' (write 1 to clear) are set by the core when an interrupt has been triggered and must be cleared by the software by writing a 1.

EIRQ<i> External interrupt <i> pending service.
Raised when the core input `external_irq[i]` is high, cleared by writing 1 to it as long as the input has been cleared too.

8.- Performance

8.1.- Synthesis Results

These are the synthesis results for the Dhrystone demo.

Table 5: **Synthesis Results for 'Dhrystone' Demo**

Target Device	Synthesis Type	F _{max}	Resources			
			CPU	Timer	UART	Total for Dhrystone demo
Cyclone-II -C7	Balanced	62 MHz	948 LEs, 29 M4Ks, 1 MUL9	88 LEs	143 LEs	1291 LEs, 29 M4Ks, 1 MUL9
Spartan-3 -5	Speed	54 MHz	(*1)			1319 LUTs, 10 BRAMs, 1 MUL18

(*1) The Spartan results include only the MCU and none of the DE-1 on-board glue logic (7-segment decoders, etc.)

The Dhrystone demo includes 12KB of ROM and 2 KB of XRAM, besides the IRAM. On the DE-1 board, it also includes some auxiliary glue logic that is uncluded in the above total count *for the Cyclone-II target only*. That glue logic accounts for 66LEs.

Results for Cyclone-II are the the actual synthesis results obtained for the Dhrystone demo on Terasic's DE-1 board (using top file [c2sb_demo.vhdl](#)).

Results for the Spartan-3 core are more speculative because the core has not yet been tested in Xilinx hardware. They include the MCU entity only and not the top entity glue logic.

The above results have been obtained with Quartus-2 11.1 sp2 and Xilinx ISE 9.2i. The synthesis has been performed unconstrained and the results must be considered illustrative only.

8.2.- Dhrystone 2.1 Benchmark

The MCU has executed a version of the Dhrystone 2.1 benchmark, [adapted for MCUs by ECROS Technology](#) and slightly modified to suit the light52 core. It has been compiled with [SDCC](#).

The benchmark has been executed on a DE-1 development board with a Cyclone-II FPGA clocked at 50 MHz using the top module entity `/vhdl/demos/c2sb/c2sb_demo.vhdl`. The benchmark executes 25000 iterations over the Dhrystone loop and produces the following results:

Dhrystone 2.1 Benchmark Results	
1641	Dhrystones per second @ 50 MHz
0.934	Dhrystone MIPS (*1)
0.0187	Dhrystone MIPS per MHz

(*1) 1 Dhrystone MIPS = 1757 Dhrystones per second

In order to give some context for this benchmark, the following table compares the results for a few representative cores:

CPU	D-MIPS	Advantage vs. light52	F _{max}
Light52	0.0187		62 MHz
Intel MCS51	0.0094	x0.5	12 MHz
CAST R8051XC-2 AF	0.0883	x4.7	35 MHz

(*1) F_{max} on a Cyclone-II FPGA

As can be seen below, this core is about twice as fast as a 12-clocker and therefore can be characterized as a 6-clocker, even if the clock count per instruction is not linearly scalable from the original.

The single-clocker [CAST's R8051XC-2 AF](#) has been selected because its feature set is not far ahead of light52's.

Note that the area performance of light52 is way better than the selected commercial core, which is amongst the smallest 8051 commercial cores. A detailed comparison would need to account for the many optional modules available in any commercial core (OCD, dual DPTR, etc.), of which R8051XC-2 AF has few.